


# Statistical Learning


Juan C. Laria

2018/11/14

... with 


# Getting Started with $k$ Nearest Neighbors





# k-NN decision boundary

k = 5 (training set)



**How does it work?**

# Euclidean distance


$$d(\mathbf{x}_1, \mathbf{x}_2) = ((\mathbf{x}_1 - \mathbf{x}_2)'(\mathbf{x}_1 - \mathbf{x}_2))^{1/2} = \sqrt{\sum_{j=1}^p (x_{1j} - x_{2j})^2}.$$




# What does it do?

$$k = 5$$


# What does it do?




# What does it do?




# What does it do?




# What does it do?




# What does it do?



# What does it do?



# What does it do?





# Practice time!




 <https://jlaria.github.io/SUsl/knn>

 [https://raw.githubusercontent.com/jlaria/SUsl/master/source/knn\\_script.R](https://raw.githubusercontent.com/jlaria/SUsl/master/source/knn_script.R)

Lunch time!



# Logistic Regression



# The intuition

We can think of the action of buying the product as a random variable defined as,

$$Y = \begin{cases} 0, & \text{Not buying} \\ 1, & \text{buying} \end{cases}$$

# The intuition

We can think of the action of buying the product as a random variable defined as,

$$Y = \begin{cases} 0, & \text{Not buying} \\ 1, & \text{buying} \end{cases}$$

Conditioning  $Y$  on the value of  $X$ , we can model the following probabilities.

$$P(Y = 1|X = x) = p(x),$$

$$P(Y = 0|X = x) = 1 - p(x).$$

# The intuition

We can think of the action of buying the product as a random variable defined as,


$$Y = \begin{cases} 0, & \text{Not buying} \\ 1, & \text{buying} \end{cases}$$

Conditioning  $Y$  on the value of  $X$ , we can model the following probabilities.

$$P(Y = 1|X = x) = p(x),$$

$$P(Y = 0|X = x) = 1 - p(x).$$

This means that  $Y|X = x$  follows a Bernoulli distribution with probability of success  $p(x)$ . Instead of modelling the response  $Y$ , we model  $p(x)$ .





# The intuition

- If we wanted to work with linear functions, say  $\beta_0 + \beta X$ , it shouldn't be  $p(x)$  the function to be approximated. Notice that,
  - $0 \leq p \leq 1$

# The intuition

- If we wanted to work with linear functions, say  $\beta_0 + \beta X$ , it shouldn't be  $p(x)$  the function to be approximated. Notice that,
  - $0 \leq p \leq 1$
- Then, the *odds*, defined as  $p/(1 - p)$  are such that,
  - $0 \leq p/(1 - p) \leq +\infty$

# The intuition

- If we wanted to work with linear functions, say  $\beta_0 + \beta X$ , it shouldn't be  $p(x)$  the function to be approximated. Notice that,

- $0 \leq p \leq 1$

- Then, the *odds*, defined as  $p/(1 - p)$  are such that,

- $0 \leq p/(1 - p) \leq +\infty$

- The odds can be approximated better by a linear function, but notice that, for  $p = 0.5$ ,  $p/(1 - p) = 1$ . This means that for  $x$  such that  $p(x) < 0.5$ , the odds are between 0 and 1, and for  $x$  such that  $p(x) > 0.5$ , the odds are between 1 and  $\infty$ .

# The intuition

- If we wanted to work with linear functions, say  $\beta_0 + \beta X$ , it shouldn't be  $p(x)$  the function to be approximated. Notice that,
  - $0 \leq p \leq 1$
- Then, the *odds*, defined as  $p/(1 - p)$  are such that,
  - $0 \leq p/(1 - p) \leq +\infty$
- The odds can be approximated better by a linear function, but notice that, for  $p = 0.5$ ,  $p/(1 - p) = 1$ . This means that for  $x$  such that  $p(x) < 0.5$ , the odds are between 0 and 1, and for  $x$  such that  $p(x) > 0.5$ , the odds are between 1 and  $\infty$ .
- To tackle this unbalance, we can take logarithms.
  - $-\infty \leq \log(p/(1 - p)) \leq +\infty$


# The intuition


- The *log-odds* may be the perfect candidates for a linear approximation.


$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta X$$

- Taking  $p$  out,

$$p = [1 + \exp(-\beta_0 - \beta X)]^{-1}$$









# How does it work?

Let  $y_1, y_2, \dots, y_N$  be the responses, and  $p_1, p_2, \dots, p_N$  the underlying probabilities that generated them from the Bernoulli model. Then, the likelihood is,

$$L(y_1 \dots y_N | p_1 \dots p_N) = \prod_{i=1}^N p_i^{y_i} (1 - p_i)^{(1-y_i)} = \prod_{i=1}^N \left( \frac{p_i}{1 - p_i} \right)^{y_i} (1 - p_i)$$

Then, the log-likelihood becomes,

$$l(y_1 \dots y_N | p_1 \dots p_N) = \sum_{i=1}^N \left\{ y_i \log \left( \frac{p_i}{1 - p_i} \right) + \log(1 - p_i) \right\}$$

Let  $\eta_i = \mathbf{x}_i^T \beta$ , and notice that the odds can be written in terms of  $\eta_i$  as

$$\frac{p_i}{1 - p_i} = \exp(\eta_i)$$

# How does it work?

Then, the log-likelihood can be written

$$l(y_1 \dots y_N | p_1 \dots p_N) = \sum_{i=1}^N \{y_i \eta_i + \log([1 + \exp(\eta_i)]^{-1})\}$$

$$l(y_1 \dots y_N | p_1 \dots p_N) = \sum_{i=1}^N \{y_i \eta_i - \log(1 + e^{\eta_i})\}$$

Therefore, the objective of logistic regression is to minimize with respect to  $\beta$  the function

$$R(\beta) = \sum_{i=1}^N \log(1 + \exp(\mathbf{x}_i^T \beta)) - \sum_{i=1}^N y_i \mathbf{x}_i^T \beta$$

# Practice time!



 <https://jlaria.github.io/SUsl/logit>

 [https://raw.githubusercontent.com/jlaria/SUsl/master/source/logit\\_script.R](https://raw.githubusercontent.com/jlaria/SUsl/master/source/logit_script.R)



# Unsupervised statistical learning

k-means clustering

Hierarchical clustering

# The intuition

<https://www.naftaliharris.com/blog/visualizing-k-means-clustering/>

# Practice time!



 <https://jlaria.github.io/SUsl/kmeans>

 <https://jlaria.github.io/SUsl/hclust>

# Thanks!

Slides created via the R package **xaringan**.

**remark.js**, **knitr**, and R Markdown.